

Comparative Clustering of Shape and S-Set Datasets:

S-sets:

Synthetic 2-d data with $N=5000$ vectors and $k=15$ Gaussian clusters with different degree of cluster overlap

P. Fränti and O. Virtajoki, "Iterative shrinking method for clustering problems", Pattern Recognition, 39 (5), 761-765, May 20

Shape Sets:

Pathbased

$N=300$, $k=3$, $D=2$ Pathbased: txt H. Chang and D.Y. Yeung, Robust path-based spectral clustering. Pattern Recognition, 2008. 41(1): p. 191-2

Spiral

$N=312$, $k=3$, $D=2$ Spiral: txt H. Chang and D.Y. Yeung, Robust path-based spectral clustering. Pattern Recognition, 2008. 41(1): p. 191-2

Jain

$N=373$, $k=2$, $D=2$ Jain: txt A. Jain and M. Law, Data clustering: A user's dilemma. Lecture Notes in Computer Science, 2005. 3776: p. 1-

Flame

N=240, k=2, D=2 Flame: txt L. Fu and E. Medico, FLAME, a novel fuzzy clustering method for the analysis of DNA microarray data. BMC bioinformatics, 2007. 8(1): p. 3.10.03.03.06

Importing modules and datasets

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
S1 = pd.read_csv("s1.txt",sep="\s+",header=None)
S4 = pd.read_csv("s4.txt",sep="\s+",header=None)
```

```
#Third row is label
pathbased = pd.read_csv("pathbased.txt",sep="\s+",header=None)
spiral = pd.read_csv("spiral.txt",sep="\s+",header=None)
jain = pd.read_csv("jain.txt",sep="\s+",header=None)
flame = pd.read_csv("flame.txt",sep="\s+",header=None)
spiral
```

	0	1	2
0	31.95	7.95	3
1	31.15	7.30	3
2	30.45	6.65	3
3	29.70	6.00	3
4	28.90	5.55	3
...
307	15.75	13.85	2
308	15.65	14.05	2
309	15.65	14.25	2
310	15.65	14.50	2
311	15.65	14.60	2

```
dfs= {"S1": [S1,15],
      "S4": [S4,15],
```

```

        "pathbased": [pathbased, 3],
        "spiral": [spiral, 3],
        "jain": [jain, 2],
        "flame": [flame, 2]

    }
    for df_name in dfs.keys():
        if df_name=="S1" or df_name=="S4":
            pass
        else:
            df=dfs[df_name][0]
            df.rename(columns={2:"true_cluster"}, inplace=True)
            df['true_cluster']=df['true_cluster']-1 # Changing the labeling from 1...k to 0..

    flame.true_cluster.unique()

```

```
array([0, 1], dtype=int64)
```

Visualization







```

plt.style.use('ggplot')

fig, axes = plt.subplots(6, 1, figsize=(5, 20))
axes = axes.ravel()
i = 0
for df in dfs.keys():
    axes[i].set_title(f'{df}')

    sns.scatterplot(x=0, y=1, data=dfs[df][0], ax=axes[i], color="r", s=10)
    i += 1
fig.tight_layout()
# plt.show()

```



```
flame
```

	0	1	true_cluster
0	1.85	27.80	0
1	1.35	26.65	0
2	1.40	23.25	1
3	0.85	23.05	1
4	0.50	22.35	1
...
235	7.50	26.20	0
236	7.50	25.65	0
237	7.05	25.85	0
238	6.90	27.15	0
239	6.15	26.90	0

```
S1.columns
```

```
Index([0, 1], dtype='int64')
```

Kmeans

Kmeans tries to minimize within cluster variability. Within cluster variability is defined as sum of squared distance of each pair of data in a cluster. The Distance is defined as Euclidean Distance. The Optimization of this objective is not trivial. Instead in the K-means algorithm a local optimum is obtained by the following operation.

Initialization:

Samples are randomly assigned to a cluster First Step:

Cluster Centroid is calculated (a vector with k dimensions each feature with the mean of samples in the cluster)

Second Step: The observations that have the minimum distance to a centroid are assigned to that cluster

First and Second Steps Are iterated .

```
from sklearn.cluster import KMeans

fig,axs=plt.subplots(6,1,figsize=(5,20))
```







```

axs=axs.ravel()
i=0
for df_name in dfs.keys():
    dataset=dfs[df_name][0]
    k = dfs[df_name][1]

    kmeans=KMeans(n_clusters=k,n_init='auto')
    kmeans.fit(dataset.loc[:,0:1])
    dataset["kmeans_cluster"]=kmeans.labels_
    axs[i].set_title(f'{df_name}')

    #sns.scatterplot(x=0,y=1,data=dataset,ax=axs[i],hue='kmeans_cluster',s=10)
    axs[i].scatter(x=dataset.loc[:,0],y=dataset.loc[:,1],c=dataset.loc[:, 'kmeans_cluster'])
    i+=1
fig.tight_layout()
#plt.show()

```



Hierarchical Clustering

Two Dissimilarity measures need to be defined here: Metric : is a measure of distance of two observations. could be euclidean,l1,l2, correlation based and etc. Linkage: is about how we find the dissimilarity between two groups of observations. Ward: minimizes the variance of clusters Single: calculates pairwise distances of observations of two clusters and considers the smallest one Complete: calculates pairwise distances of observations of two clusters and considers the largest one Average: calculates pairwise distances of observations of two clusters and averages them Centroid: Dissimilarity between center of one cluster and the center of the other cluster. Center is a mean feature vector of dim p

Documentation of Scikit learn

Metric : str or callable, default=None Metric used to compute the linkage. Can be “euclidean”, “l1”, “l2”, “manhattan”, “cosine”, or “precomputed”. If set to None then “euclidean” is used. If linkage is “ward”, only “euclidean” is accepted. If “precomputed”, a distance matrix is needed as input for the fit method.

linkage : {'ward', 'complete', 'average', 'single'}, default='ward' Which linkage criterion to use. The linkage criterion determines which distance to use between sets of observation. The algorithm will merge the pairs of cluster that minimize this criterion.

- ward minimizes the variance of the clusters being merged.
- average uses the average of the distances of each observation of the two sets.
- complete or maximum linkage uses the maximum distances between all observations of the two sets.
- single uses the minimum of the distances between all observations of the two sets

Average Linkage

```
from sklearn.cluster import AgglomerativeClustering







fig,axs=plt.subplots(6,1,figsize=(5,20))
axs=axs.ravel()
i=0
for df_name in dfs.keys():
    dataset=dfs[df_name][0]
    k = dfs[df_name][1]

    hierarchical_cluster = AgglomerativeClustering(n_clusters=k, metric='euclidean', linka
    labels = hierarchical_cluster.fit_predict(dataset.loc[:,0:1])
```



```
dataset["avg_hcluster"] = labels
axs[i].set_title(f'{df_name}')

#sns.scatterplot(x=0,y=1,data=dataset,ax=axs[i],hue='kmeans_cluster',s=10)
axs[i].scatter(x=dataset.loc[:,0],y=dataset.loc[:,1],c=labels,cmap='rainbow',s=5)
i+=1
fig.tight_layout()
#plt.show()
```









Single Linkage

```
from sklearn.cluster import AgglomerativeClustering

fig,axs=plt.subplots(6,1,figsize=(5,20))
axs=axs.ravel()
i=0
for df_name in dfs.keys():
    dataset=dfs[df_name][0]
    k = dfs[df_name][1]

    hierarchical_cluster = AgglomerativeClustering(n_clusters=k, metric='euclidean', linka
    labels = hierarchical_cluster.fit_predict(dataset.loc[:,0:1])
    dataset["single_hcluster"]= labels
    axs[i].set_title(f'{df_name}')

    #sns.scatterplot(x=0,y=1,data=dataset,ax=axs[i],hue='kmeans_cluster',s=10)
    axs[i].scatter(x=dataset.loc[:,0],y=dataset.loc[:,1],c=labels,cmap='rainbow',s=5)
    i+=1
fig.tight_layout()
#plt.show()
```









Complete Linkage

```
from sklearn.cluster import AgglomerativeClustering

fig,axs=plt.subplots(6,1,figsize=(5,20))
axs=axs.ravel()
i=0
for df_name in dfs.keys():
    dataset=dfs[df_name][0]
    k = dfs[df_name][1]

    hierarchical_cluster = AgglomerativeClustering(n_clusters=k, metric='euclidean', linka
    labels = hierarchical_cluster.fit_predict(dataset.loc[:,0:1])
    dataset["complete_hcluster"]= labels
    axs[i].set_title(f'{df_name}')

    #sns.scatterplot(x=0,y=1,data=dataset,ax=axs[i],hue='kmeans_cluster',s=10)
    axs[i].scatter(x=dataset.loc[:,0],y=dataset.loc[:,1],c=labels,cmap='rainbow',s=5)
    i+=1
fig.tight_layout()
#plt.show()
```



Spectral Clustering

For this algorithm a similarity graph must be constructed. In this graph the edges represent the similarity between nodes (objects). Similar nodes are connected and the edge weight represents the similarity and if the similarity between two nodes is not above a certain threshold, they don't get connected. In this case the problem boils down to a graph partition problem, where we have to get the connected components of the graph such that edges in each component have high weight and edges between components have low weight. Two metrics should be defined here: 1-The measure of similarity between nodes: It's usually defined as $s(x_i, x_j) = \exp(-(x_i - x_j)^2 / c)$

2- How we construct the graph similarity :- Mutual K-nearest-neighbor graph: If node v_i and v_j are both k-nearest neighbors of each other, they are connected and $w_{ij} = s_{ij}$ - Fully connected graph: connect all points and weight each edge with s_{ij} - epsilon neighborhood graph

affinity : str or callable, default='rbf' How to construct the affinity matrix. - 'nearest_neighbors': construct the affinity matrix by computing a graph of nearest neighbors. - 'rbf': construct the affinity matrix using a radial basis function (RBF) kernel.

```
from sklearn.cluster import SpectralClustering
spectral_cluster = SpectralClustering(n_clusters=k, affinity="nearest_neighbors")







from sklearn.cluster import SpectralClustering

fig,axs=plt.subplots(6,1,figsize=(5,20))
axs=axs.ravel()
i=0
for df_name in dfs.keys():
    dataset=dfs[df_name][0]
    k = dfs[df_name][1]

    spectral_cluster = SpectralClustering(n_clusters=k, affinity="nearest_neighbors")
    spectral_cluster.fit(dataset.loc[:,0:1])
    dataset["spectral_cluster"]= spectral_cluster.labels_
    axs[i].set_title(f'{df_name}')

    #sns.scatterplot(x=0,y=1,data=dataset,ax=axs[i],hue='kmeans_cluster',s=10)
    axs[i].scatter(x=dataset.loc[:,0],y=dataset.loc[:,1],c=dataset["spectral_cluster"],cmap=cm)
    i+=1
fig.tight_layout()
plt.show()
```

```
D:\Python3.11\Lib\site-packages\sklearn\manifold\_spectral_embedding.py:273: UserWarning: Gr
warnings.warn(
```

```

for df_name in dfs.keys():

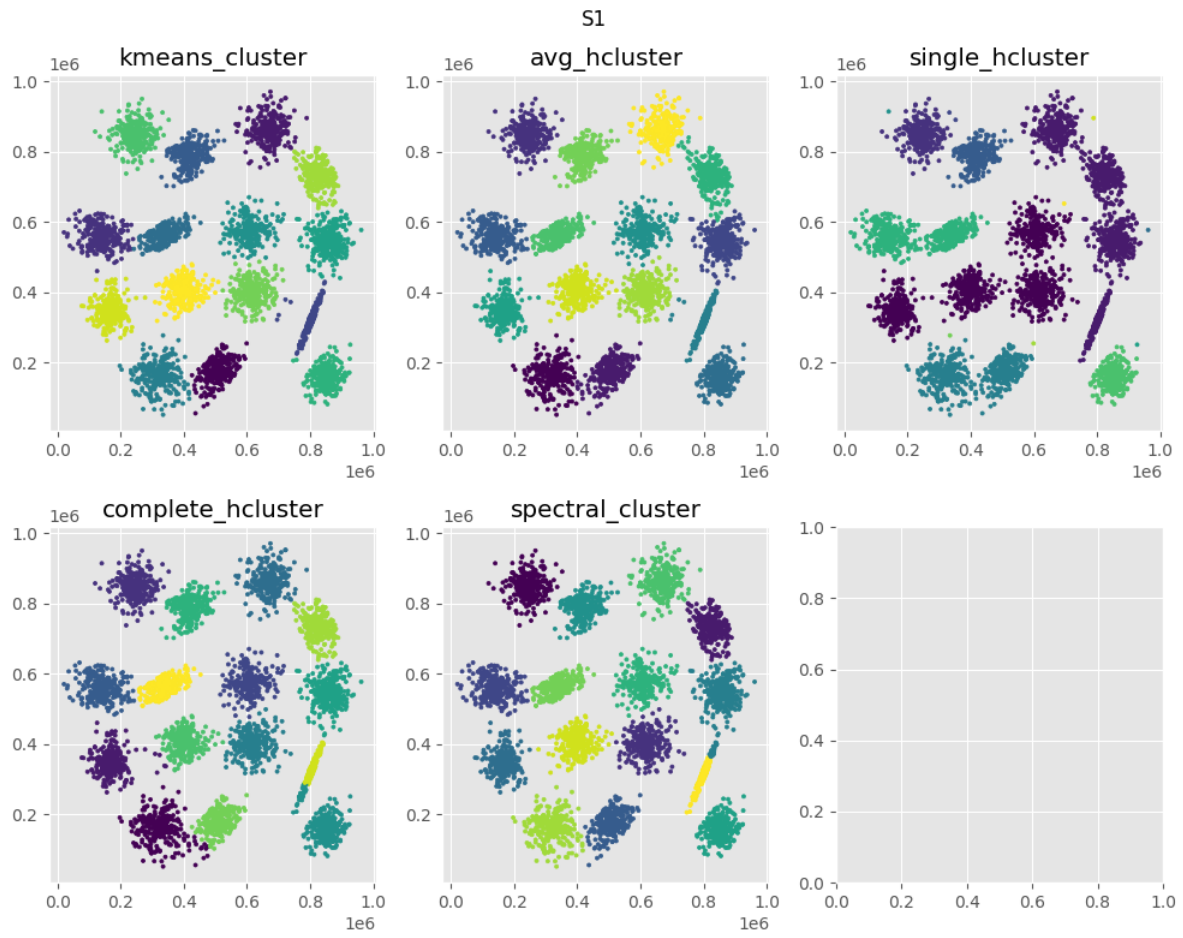
    fig,axs=plt.subplots(2,3,figsize=(10,8))
    axs=axs.ravel()
    dataset=dfs[df_name][0]
    k = dfs[df_name][1]
    fig.suptitle(f"{df_name}")
    for i in range(6):
        try:

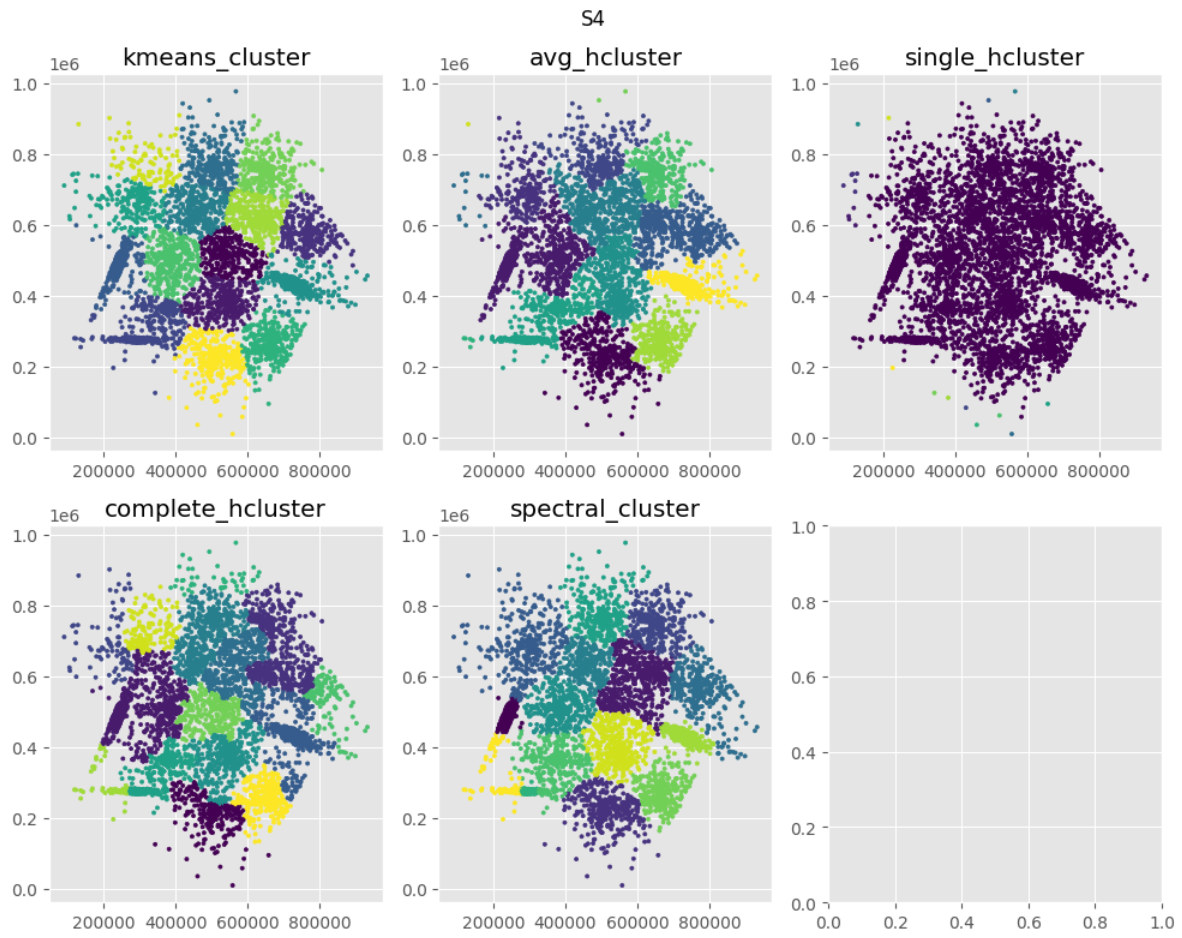
            axs[i].set_title(f'{dataset.iloc[:,i+2].name}')

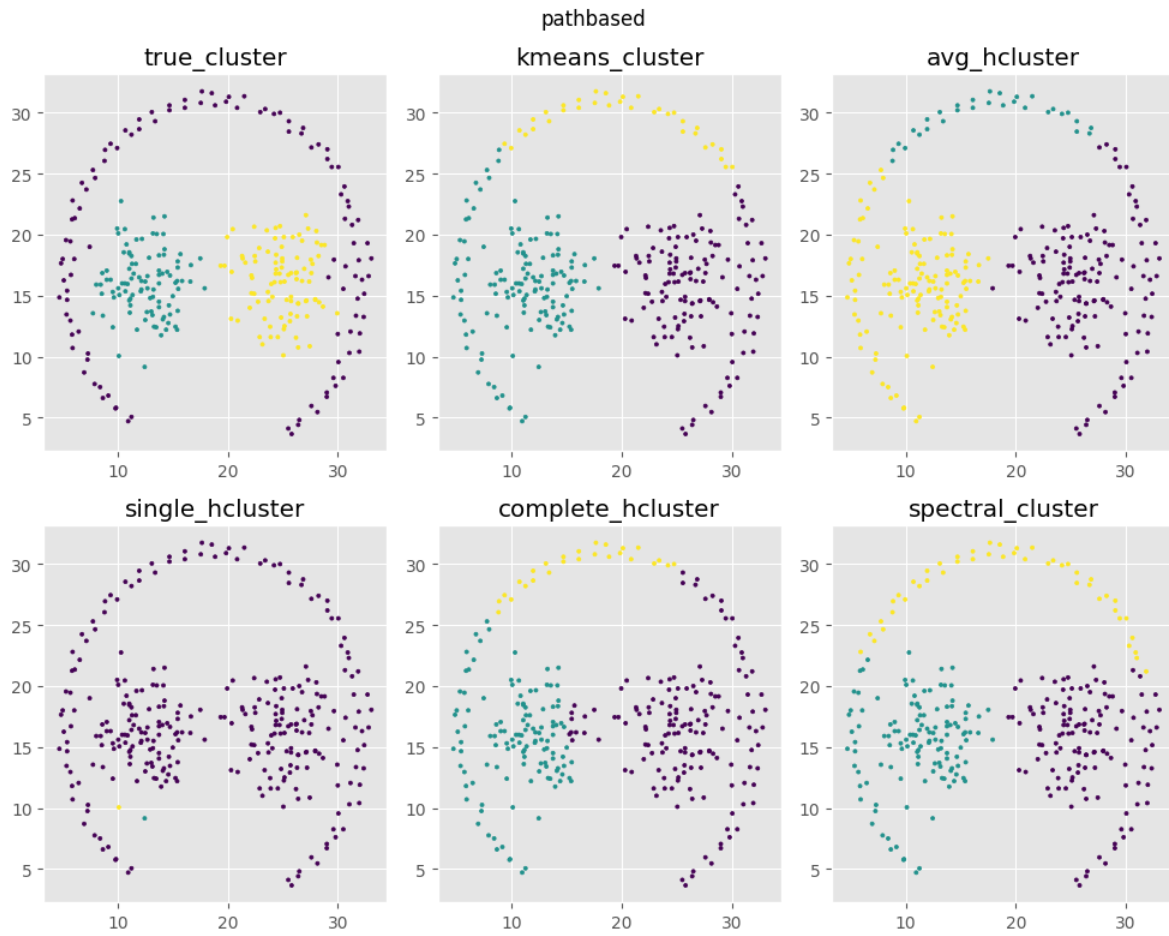
            #sns.scatterplot(x=0,y=1,data=dataset,ax=axs[i],hue='kmeans_cluster',s=10)
            #print(f"x:{x}\ny:{y}")
            axs[i].scatter(x=dataset.loc[:,0],y=dataset.loc[:,1],c=dataset.iloc[:,i+2],cma
        except:
            pass
    fig.tight_layout()

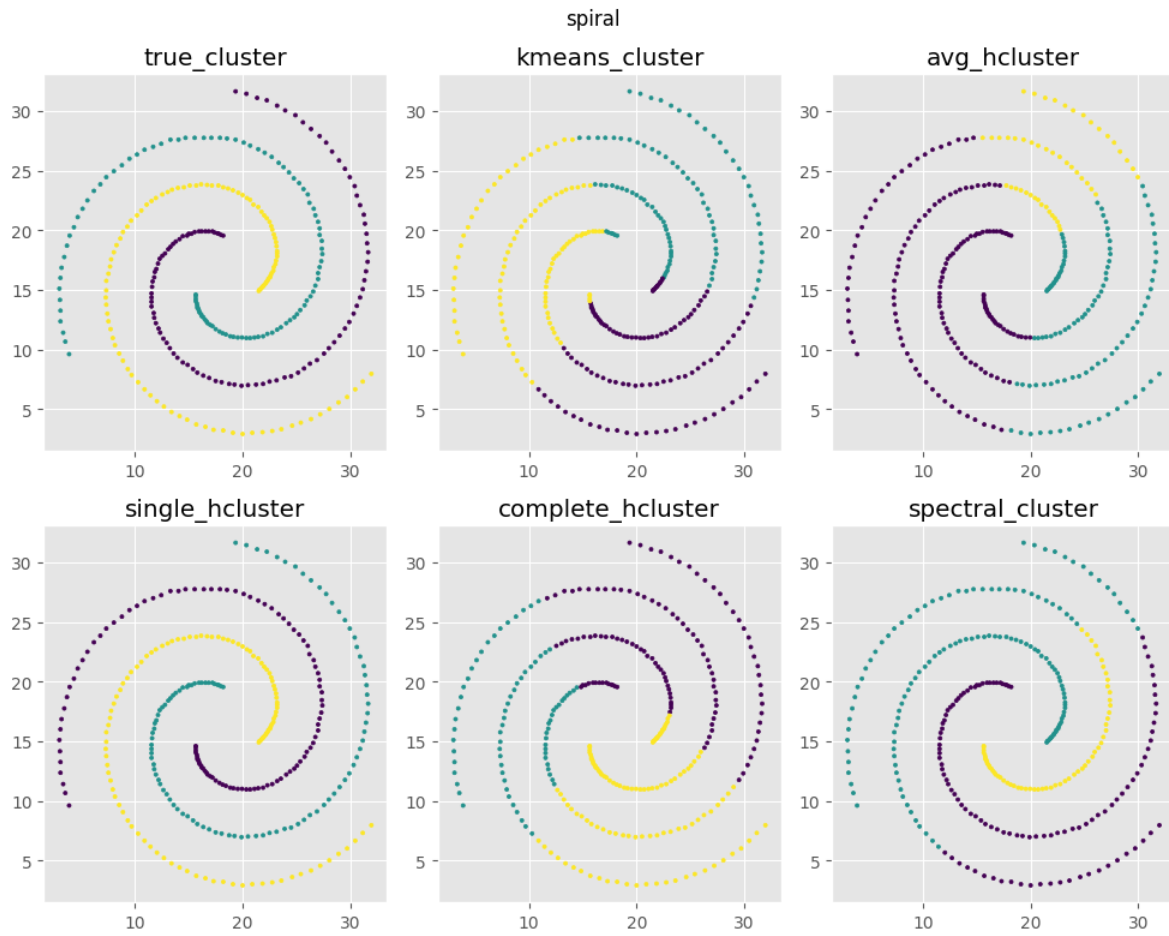
    #plt.show()


```














Purity index :

```
from sklearn.metrics import confusion_matrix, classification_report
for df_name in dfs.keys():
    if df_name=="S1" or df_name=="S4":
        continue
    dataset=dfs[df_name][0]
    k = dfs[df_name][1]
    N= len(dataset)
    true= dataset.true_cluster
    #print()
    plt.suptitle(f"{df_name}")
    fig,axs=plt.subplots(2,3,figsize=(10,8))
```


- **Purity:** Purity is a measure of the extent to which clusters contain a single class.^[37]

Its calculation can be thought of as follows: For each cluster, count the number of data points from the most common class in said cluster. Now take the sum over all clusters and divide by the total number of data points. Formally, given some set of clusters M and some set of classes D , both partitioning N data points, purity can be defined as:

$$\frac{1}{N} \sum_{m \in M} \max_{d \in D} |m \cap d|$$

Figure 1: image.png

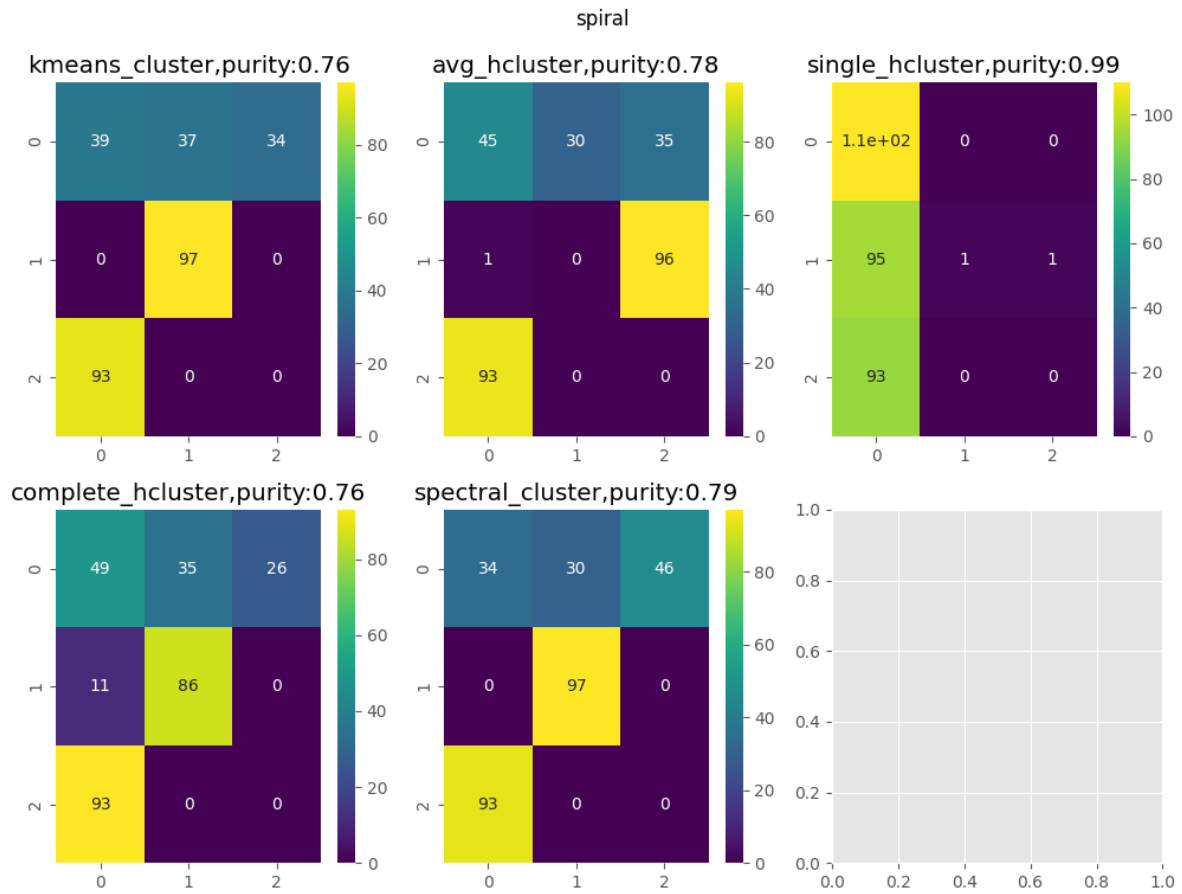
```

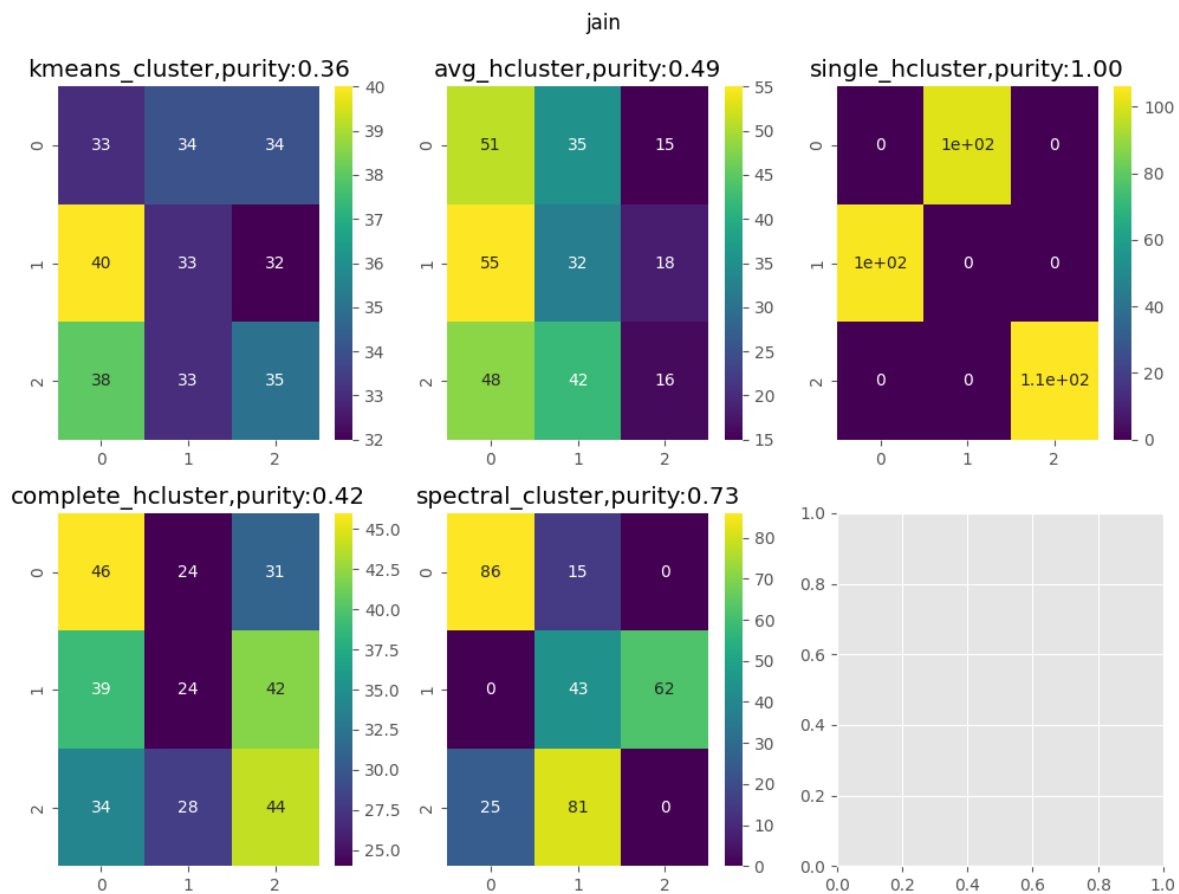
axs=axs.ravel()
for i in range(5):
    predicted=dataset.iloc[:,i+3]
    confusion_mat=confusion_matrix(true, predicted)
    #print(confusion_mat)
    purity= np.sum(np.amax(confusion_mat,axis=1))/len(dataset)
    #print(f"Purity Index of {dataset.iloc[:,i+3].name}:{purity}")
    axs[i].set_title(f"{dataset.iloc[:,i+3].name},purity:{purity:.2f}")

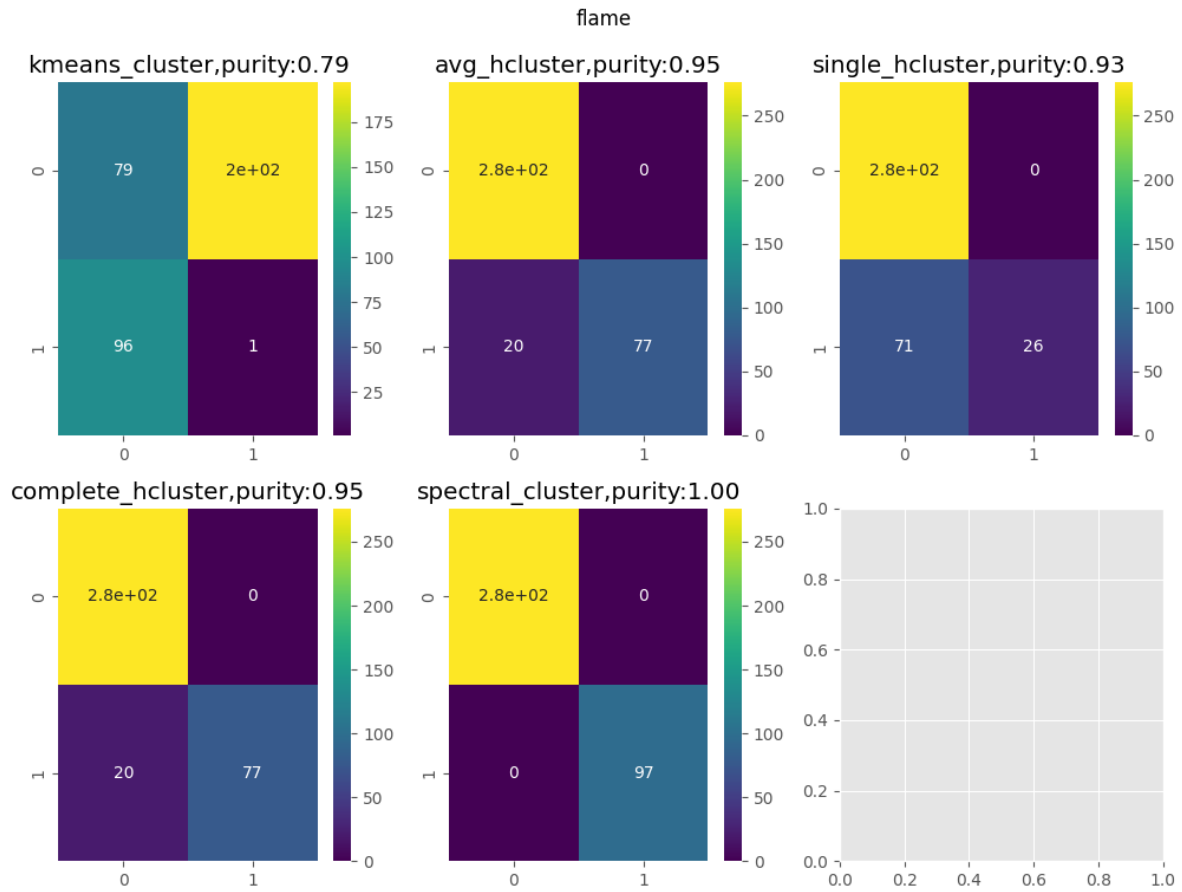
    sns.heatmap(confusion_mat,ax=axs[i],annot=True,cmap="viridis")
    fig.tight_layout(rect=[0, 0.03, 1, 0.95])


```

<Figure size 640x480 with 0 Axes>









Rand index :

Rand index [\[edit\]](#)

(right) algorithms. The calculated Adjusted Rand index for these two clusterings is $ARI \approx 0.94$

Definition [\[edit\]](#)

Given a [set](#) of n elements $S = \{o_1, \dots, o_n\}$ and two [partitions](#) of S to compare,

$X = \{X_1, \dots, X_r\}$, a partition of S into r subsets, and $Y = \{Y_1, \dots, Y_s\}$, a partition of S into s subsets, define the following:

- a , the number of pairs of elements in S that are in the **same** subset in X and in the **same** subset in Y
- b , the number of pairs of elements in S that are in **different** subsets in X and in **different** subsets in Y
- c , the number of pairs of elements in S that are in the **same** subset in X and in **different** subsets in Y
- d , the number of pairs of elements in S that are in **different** subsets in X and in the **same** subset in Y

The Rand index, R , is:^{[1][2]}

$$R = \frac{a + b}{a + b + c + d} = \frac{a + b}{\binom{n}{2}}$$

Intuitively, $a + b$ can be considered as the number of agreements between X and Y and $c + d$ as the number of disagreements between X and Y .

Since the denominator is the total number of pairs, the Rand index represents the *frequency of occurrence* of agreements over the total pairs, or the probability that X and Y will agree on a randomly chosen pair.

$\binom{n}{2}$ is calculated as $n(n-1)/2$.

Similarly, one can also view the Rand index as a measure of the percentage of correct decisions made by the algorithm. It can be computed using the following formula:

Similarly, one can also view the Rand index as a measure of the percentage of correct decisions made by the algorithm. It can be computed using the following formula:


$$RI = \frac{TP + TN}{TP + FP + FN + TN}$$


where TP is the number of true positives, TN is the number of [true negatives](#), FP is the number of [false positives](#), and FN is the number of [false negatives](#).

```
from sklearn.metrics.cluster import rand_score
methods=S1.columns[2:]
for df_name in dfs.keys():
    dataset=dfs[df_name][0]
    k = dfs[df_name][1]
    N= len(dataset)
    fig,axs=plt.subplots(1,1,figsize=(10,8))
    randind_mat=np.zeros((5,5))
    for i in range(5):
        for j in range(5):
            randind_mat[i,j]= rand_score(dataset.loc[:,methods[i]],dataset.loc[:,methods[j]])
```

```
sns.heatmap(randind_mat,annot=True,cmap="viridis",xticklabels=methods,yticklabels=methods)
fig.suptitle(f"{df_name}")


fig.tight_layout(rect=[0, 0.03, 1, 0.95])
```








spiral



jain





References

Websites:

- 1- https://www.w3schools.com/python/python_ml_k-means.asp
- 2- https://www.w3schools.com/python/python_ml_hierarchial_clustering.asp
- 3- <https://scikit-learn.org/stable/modules/generated/sklearn.cluster.SpectralClustering.html>
- 4- https://scikit-learn.org/stable/modules/generated/sklearn.metrics.confusion_matrix.html
- 5- <https://stats.stackexchange.com/questions/95731/how-to-calculate-purity>